

CS 143

Warren Kim

Contents

1	Overview	3
1.1	Purpose of a Database	3
1.2	Abstraction Layers	3
1.3	Instances and Schema	3
1.4	Data Models	3
1.4.1	Relational	4
1.4.2	Entity-Relationship (ER)	4
1.4.3	Object-Oriented	4
1.4.4	Document (Semi-Structured)	4
1.4.5	Network/Hierarchical/Graphical	4
1.4.6	Vector	4
1.4.7	Key-Value	4
1.5	Database Languages	4
1.5.1	Data Manipulation Language	5
1.5.2	Data Definition Language	5
1.6	Data Storage and Querying	5
1.7	Defining a Schema	5
2	Keys	6
2.1	Superkey	6
2.2	Candidate Key	6
2.3	Primary Key	6
2.4	Foreign Key	7
3	Relational Algebra	8
3.1	Selection	8
3.2	Projection	8
3.3	Cartesian Product	9
3.4	Aggregation	9
3.5	Rename	10
3.6	Set Operations	10
3.7	Order of Precedence	10
4	Joins	11
4.1	Natural Join	11
4.2	Theta Join	11
4.3	Inner Joins	12
5	PostgreSQL Data Types	13
5.1	Numbers	13
5.2	Strings	13
5.3	Binary Data and Booleans	13
5.3.1	Binary Data	13
5.3.2	Boolean	14
5.4	Date/Time	14

5.5	02/12/24	15
5.5.1	From RDBS To Other Systems	15
5.5.2	OLAP and Data Warehousing	15

Overview

1.1 Purpose of a Database

We will be studying (mostly) Relational DataBase Management Systems (RDBMS).

Definition: Database

A **database** abstracts how data is stored, maintained, and processed. It is a system that uses advanced data structures to store and index data.

A database abstracts away the data integrity and file management aspect of CRUD operations. Moreover, a database provides us with a single location for all of the data, even if the database itself is distributed.

1.2 Abstraction Layers

There are three layers of abstraction: physical, logical, and view.

Definition: Physical Abstraction

The **physical abstraction** defines the data and its relationships to other data within the database.

Definition: Logical Abstraction

The **logical abstraction** deals with how we interface with the database.

Definition: View Abstraction

The **view abstraction** refers to specific use cases and filters the data from the logical abstraction.

We start by learning the logical abstraction.

1.3 Instances and Schema

Definition: Schema and Instance

A **schema**^a is the overall design of a database. It defines the structure of the data as well as how it is organized.

An **instance** of a database is the actual set of data stored in the database at a particular moment in time.

^aNote: schema can also refer to a relation (table).

1.4 Data Models

Data models define how we design databases and interact with data. We want to answer the following:

- How do we define data?
- How do we encode relationships among data?
- How do we impose constraints on data?

Data models are either an Implementation model or a Design mechanism. Implementation models build databases from the ground up while design mechanisms are implemented as features in a database. We discuss five major types (an several niche ones).

1.4.1 Relational

In a relational model, all data is stored as a *relation*¹. Rows represent individual n -tuple units (*records*). Columns represent (typed) *attributes* common to all records in the relations.

1.4.2 Entity-Relationship (ER)

An entity-relationship model uses a collection of basic objects (*entities*) and define *relationships* among them.

1.4.3 Object-Oriented

The object-oriented model is similar to OOP with encapsulation, methods, and object identity. It was originally an implementation model but is now a design mechanism.

1.4.4 Document (Semi-Structured)

A document model stores records as *documents*, which do *not* have an enforced schema. This allows for more versatility in the type of data stored in the database.

1.4.5 Network/Hierarchical/Graphical

A graph model is analogous to how we think. Records are stored as *nodes* and relationships between records as *edges*.

1.4.6 Vector

A vector model stores records as *vectors* in \mathbb{R}^n , and are stored in a way that enables efficient retrieval and comparison (e.g. nearest neighbor[s]).

1.4.7 Key-Value

A key-value model stores data as a key-value pair (typically using a hash function). In this model, data typically lives in RAM as opposed to disk.

1.5 Database Languages

There are two main semantic systems when working with databases:

- Data Manipulation Language (DML)
- Data Definition Language (DDL)

Note that a relational model typically uses SQL for both DDL and DML.

¹Note: tables are an implementation of relations.

1.5.1 Data Manipulation Language

DML's can either be procedural or declarative.

Definition: Query

A **query** is a written expression to retrieve or manipulate data.

Aside: A Note on SQL

SQL is a declarative language, and as such, it is hard to perform sequential or nontrivial^a computations in SQL. To remedy this, a common option is to write an **ETL job** in another language (pick one). We **E**xtract the data from the database (using a connection driver), **T**ransform the data using another language (pick one!), and **L**oad the data into a new table using the same driver. We can schedule these jobs using something like **cron**.

^aNontrivial: Any computation where we have to specify *how* to perform the computation.

1.5.2 Data Definition Language

DDL's specify a schema: a collection of attribute names and data types, consistency constraints, and optionally storage structure and access methods. There are four types of consistency constraints:

- Domain constraints define the domain of an attribute (e.g. `tinyint`, `enum`, etc.).
- Assertions are business rules that must hold true (e.g. an enforced prerequisite for a class must be present in your transcript before you can add a class to your study list).
- Authorization determines who can do what (e.g. full CRUD, read-only, etc.).
- Referential integrity ensures that links from one table to another must be defined (Suppose we have two relations R, R' . If there is a link $f : R \rightarrow R'$, then f is surjective).

1.6 Data Storage and Querying

Definition: Storage Manager

A **storage manager** that abstracts away how the data is laid out on disk.

A storage manager is helpful because reading data from disk to RAM is *slow*, and the storage manager handles swapping² and makes retrieval efficient.

Definition: Query Manager

A **query manager** takes the DML statements and organize them into a *query plan*^a that “compiles” a query (using relational algebra) and executes the instruction(s).

^aNote: The query plan dictates the performance of a query.

1.7 Defining a Schema

A schema can be written as `relation(attribute1, ..., attributen)` where underlined attributes represent the primary key.

²Swapping: Virtual memory in CS111!

Keys

Aside: A Note on Context and Instance

Based on **context** means that the given data is a subset of the complete dataset.
Based on **instance** means that we treat the given data as the complete dataset.

2.1 Superkey

Definition: Superkey

A **superkey** is a set of one or more attributes that uniquely identifies a record (tuple) and distinguishes it from all other records in the relation.

Formally, let R be a relation with a set $S = \{a_1, a_2, \dots, a_n : a \text{ is an attribute of } R\}$. A **superkey** is a subset $s \subseteq S$ such that s uniquely identifies each n -tuple in R .

The superkey $s = S = \{a_1, a_2, \dots, a_n\} = \bigcup_{i=1}^n \{a_i\}$ is called the *trivial superkey*. Additionally, \emptyset is not a superkey. Further note that for every relation R , there exists at most $2^n - 1$ superkeys where n is the number of attributes.

2.2 Candidate Key

Definition: Candidate Key

A **candidate key** is a superkey such that no subset of the candidate key is a superkey; i.e. it is the minimal superkey.

Formally, let R be a relation with a set $S = \{a_1, a_2, \dots, a_n : a \text{ is an attribute of } R\}$. A **candidate key** is a superkey $s \subseteq S$ such that for every proper subset $t \subsetneq s$, t is not a superkey.

Candidate keys may vary in length, and the attributes of a candidate key may be NULL as long as it uniquely identifies an n -tuple in the relation.

2.3 Primary Key

Definition: Primary Key and Composite Key

A **primary key** is a candidate key (chosen by the database designer) to enforce uniqueness for a particular use case.

The primary key is typically chosen to be the minimal candidate key for simplicity. The attributes of a primary key may not be NULL.

2.4 Foreign Key

Definition: Foreign Key

A **foreign key** is a set of attributes that links tuples of two relations.

Formally, let R, R' be relations with sets $S = \{a_1, a_2, \dots, a_n : a \text{ is an attribute of } R\}, S' = \{a'_1, a'_2, \dots, a'_n : a' \text{ is an attribute of } R'\}$. A **foreign key** is a key $s \subseteq S$ of R that maps to the primary key $p \subseteq S'$ of R' .

Foreign keys are used to enforce referential integrity constraints; i.e. foreign keys in a relation R are used to protect data in R from being orphaned and/or inconsistent. Given two relations R, R' related via a foreign key, R' is said to be the *referring* relation and R the *referred* relation.

Let two relations R, S be related via a foreign key, where S is the *referring* relation and R is the *referred* relation. Suppose we want to remove an n -tuple $r \in R$. Then there are two cases:

Case 1 If there is no $s \in S$ such that $s \mapsto r$, we simply remove r .

Case 2 If there is at least one $s \in S$ such that $s \mapsto r$, we can either throw an error to prevent the deletion of r or *cascade*¹ the delete.

¹Cascade: Delete r and all $s \in S$ that refer to r .

Relational Algebra

3.1 Selection

Definition: Selection

Selection retrieves a subset of tuples from a *single* relation R that satisfies some predicate ψ and returns a new relation $R' \subseteq R$, and is defined by

$$\sigma_{\psi}(R) = R' = \{t \in R : \psi(t)\}$$

where ψ is a boolean predicate on attributes and values with respect to a unary or binary operator^a

^aWe may use the following operators: $\{=, \neq, <, >, \leq, \geq, \neg\}$.

We can build complex predicates using conjunction \wedge (and) or disjunction \vee (or).

Note: that selection σ is the most analogous to WHERE in SQL.

Below are a list of examples of selection, assuming all attributes and relations are well-defined:

- (i) $\sigma_{(\text{dislikes} < \text{likes})}(\text{youtube_video})$
- (ii) $\sigma_{(\text{cat_id}=17)}(\text{youtube_video})$
- (iii) $\sigma_{([\text{dislikes} < \text{likes}] \wedge [\text{views} > 1000000] \wedge [\text{cat_id}=24])}(\text{youtube_video})$
- (iv) $\sigma_{(\text{dislikes} < \text{likes})}(\sigma_{(\text{views} > 1000000)}(\sigma_{(\text{cat_id}=24)}(\text{youtube_videos})))$

Note that (iii) and (iv) are equivalent.

3.2 Projection

Definition: Projection

Projection extracts attributes from a set of tuples and removes duplicates. Given a relation R , n -tuple t , and a set of attributes a_1, \dots, a_n ,

$$\Pi_{a_1, \dots, a_n}(R) = \{t[a_1, \dots, a_n] : t \in R\}$$

Projection is usually the last (outermost) operation done on a relation.

Aside: Projection?

We call it a projection because we are collapsing an n -tuple down to an $(n - k)$ -tuple. That is, we take the n -tuples in a relation R_n and collapse them into a set of $(n - k)$ -tuples in a new relation R'_{n-k} .

Here, R_n is a relation with n attributes.

Projections can be generalized to “create” new attributes or rename attributes using the \rightarrow notation.

Example

We can apply arbitrary expressions to existing attributes (and create another one) by doing $\Pi_{\text{likes}/(\text{likes}+\text{dislikes})\rightarrow\text{interactions}}(R)$ or rename attributes by doing $\Pi_{\text{likes}\rightarrow\text{thumbs_up}}(R)$

Example

Consider the following relation R with $\Pi_{A,B}(R)$:

A	B	C		A	B		A	B
α	β	δ	$\xrightarrow{\text{Extract}_{A,B}}$	α	β	$\xrightarrow{\Pi_{A,B}(R)}$	α	β
α	β	γ		α	β			
α	β	λ		α	β			

Note: that projection Π is the most analogous to SELECT DISTINCT in SQL.

3.3 Cartesian Product

Definition: Cartesian Product

A **Cartesian product** forms all possible pairs of tuples. Given relations R, S ,

$$R \times S = \{(r, s) : r \in R \wedge s \in S\}$$

Example

Suppose we have two relations R, S defined below:

A	B		A	B
α	β	$R :=$	α	γ
α	γ		Δ	η

Then,

A	B	C	D
α	β	α	γ
β	γ	Δ	η
Δ	η	α	γ
Δ	η	β	γ

Cartesian products are *very* expensive since they require a lot of compute power, ram, and disk space.

3.4 Aggregation

Definition: Aggregation

The aggregation operator (γ or \mathcal{G}) is a function on groups of tuples in a relation to summarize them. Common ones include: SUM, AVG, MIN, MAX, COUNT, etc. Given a relation R ,

$${}_A\gamma_F(R) = {}_A\mathcal{G}_F(R)$$

where $A := \{\text{attributes to group by}\}$, $F := \{\text{functions to apply}\}$

3.5 Rename

Definition: Rename

The rename operator (ρ) renames relations or attributes. Given a relation with name R , renaming a relation looks like $\rho_{R'}(R)$, where R' is the new name. Renaming an attribute in R looks like $\rho_{a'/a}(R)$, where a' is the new name.

Note: We must rename one of the R 's when doing $R \times R$. That is, we must have $\rho_{R'}(R) \times R$ or $R \times \rho_{R'}(R)$.

3.6 Set Operations

Definition: Set Operations (Union, Intersection, Set Difference)

Let R, S be two sets of tuples. Then, union is defined to be

$$R \cup S = \{r_1, \dots, r_{|R|}, s_1, \dots, s_{|S|} : r_i \in R \vee s_j \in S\}$$

intersection is defined as

$$R \cap S = \{t : t \in R \wedge t \in S\}$$

and set difference is defined as

$$R - S = R \setminus S = \{t : t \in R \wedge t \notin S\}$$

3.7 Order of Precedence

The order of precedence from highest to lowest is as follows:

$$(\sigma, \Pi, \rho), (\times, \bowtie), \cap, (\cup, -)$$

Joins

Definition: Join

A **join** merges tuples from two relations R, S based on some contextually related attribute(s) in both relations. The resulting relation contains tuples of the form (r, s) where $r \in R, s \in S$.

A **join key** is the set of attribute(s) that are used to join R and S . Note that a join key is *completely unrelated* to do with uniqueness.

There are two types of joins: natural and theta.

4.1 Natural Join

Definition: Natural Join

A **natural join** \bowtie is a join where the join key is determined by the RDBMS. The simplest natural join is defined using the cartesian product:

$$R \bowtie S = \Pi_{R \cup S} (\sigma_{R.k=S.k}(R \times S)) = \{(r, s) : r \in R \wedge s \in S \wedge (r[k] = s[k])\}$$

The natural join is also characterized as an *equijoin*.

Edge Cases

1. If we have two relations R, S that have no common attributes ($k = \emptyset$), then

$$R \bowtie S = \{(r, s) : r \in R \wedge s \in S \wedge (r[k] = s[k])\} = R \times S$$

because the empty set is unique.

2. If we have two relations R, S that have common attributes but no matches, then

$$R \bowtie S = \{(r, s) : r \in R \wedge s \in S \wedge (r[k] = s[k])\} = \emptyset$$

because $r[k] = s[k]$ is always false.

4.2 Theta Join

Definition: Theta Join

A **theta join** \bowtie_{θ} is a join where the join key and condition are specified. Mathematically,

$$R \bowtie_{\theta} S = \sigma_{\theta}(R_1 \times R_2) = \{(r, s) : r \in R \wedge s \in S \wedge \theta((r, s))\}$$

where θ is the join condition.

Note: For any join, if there is a name clash in either the relation or attribute(s), we must alias them.

Example

Suppose we have two relations R, S defined below:

$$R := \begin{array}{c|c} A & B \\ \hline \alpha & \beta \\ \beta & \beta \end{array}, S := \begin{array}{c|c} A & B \\ \hline \beta & \alpha \\ \alpha & \gamma \end{array}$$

Define $\theta := R.A = S.A$. Then,

$$R \bowtie_{\theta} S = R \bowtie_{R.A=S.A} S = \begin{array}{c|c|c|c} A & B & C & D \\ \hline \alpha & \beta & \alpha & \gamma \\ \beta & \beta & \beta & \alpha \end{array}$$

4.3 Inner Joins

Definition: Theta Join

An **inner join** between two relations R, S is a theta join that omits elements that do not satisfy the join condition θ .

PostgreSQL Data Types

The ANSI SQL standard defines the following data types:

- (i) numeric
- (ii) string/text
- (iii) binary
- (iv) dates and times

Aside: Promoted

It is always good practice to only promote types to increase precision, and never demote.

5.1 Numbers

Numeric data types have the following forms:

- `int(eger)` (4 bytes)
- `smallint` (2 bytes)
- `bigint` (8 bytes)
- `decimal(n, d)`, where n is the number of digits and p is the number of digits that appear after the decimal point. `numeric/decimal` is slow.
- `real, double precision`
- `float(n)`, where n is the precision.
- `NaN` is specified with quotes, and `NaN == NaN` is true and `NaN > x` for all $x \neq \text{NaN}$. Further, operations on `NaN` return `NaN`.
- `Infinity` must be quoted.

5.2 Strings

String data types have the following forms:

- `char(n)` is a fixed-length character array of length n .
- `varchar(n)` is a variable-length character array of length $\leq n$.

5.3 Binary Data and Booleans

5.3.1 Binary Data

`bytea` accepts either hex or escape format.

5.3.2 Boolean

The following are valid:

- TRUE, 't', 'true', 'y', 'yes', 'on', '1'
- FALSE, 'f', 'false', 'n', 'no', 'off', '0'

5.4 Date/Time

String data types have the following forms:

- date (YYYY-MM-DD)
- time (HH:MM:SS)
- timestamp (YY-MM-DD)

5.5 02/12/24

5.5.1 From RDBS To Other Systems

We'll come back to RDBMS in the context of indexing, transactions, and join algorithms. PostgreSQL and other RDBMS are called OLTP systems: OnLine Transaction Processing. They are designed for frequent and **random** interactive use. They are typically used in production¹.

OLTP

1. Good for simple queries (e.g. **SELECT**, basic joins)
2. Typically used for quick reads/writes that follow no real pattern.
3. Based on *transactions* or individual events, and these events are stored as rows in RDBMS.
4. Strives to minimize redundancy and requires joins to exploit relationships in data; i.e. normalization.
5. Designed for production use.
6. Abstracts data as a 2D representation (rows/columns) called a table.

5.5.2 OLAP and Data Warehousing

OnLine Analytical Processing is **not** meant to be accessed in production. The data is meant to be used by internal users. OLAP is optimized for reads with low latency. Writes are a little slower. Typically, other automated systems write into an OLAP; e.g. ETL jobs, piping, etc. We typically do batch writes and do *ad hoc* writing.

Note: OLAP/DW are completely different systems from RDBMS/OLTP. But they both use SQL. The difference is how data is stored and how operations are optimized.

OLAP

OLAP is the concept, and Data Warehouses are an implementation of them. Pure OLAP systems aren't used very often anymore, but many of the concepts are used today in relational databases or data warehouses.

We work with **cubes** in OLAP. But, flattened, it still looks like a table. e.g.

location	product	shipping	revenue
USW	Android Tablet	2017	\$5M

OLAP and DW are used for batch processing, and not for production systems. They are optimized for low latency reads of aggregated or precomputed data. Typically, we have a lot of redundancy and denormalized tables. Some use cases include:

1. Reporting (Ledgering)
2. Dashboarding

We typically write ad hoc into the system using something like an ETL job during a low usage period (e.g. overnight).

DB $\xrightarrow{\text{extract}}$ Transform $\xrightarrow{\text{load}}$ Data Warehouse.

¹Production: power and app or a business. For example, a POS system.

ETL Job

An ETL job has three steps:

1. **Extract:** Pull data from source, usually an RDBMS.
2. **Transform:** Perform transformation on data.
3. **Load:** Write the raw data into the data warehouse. Under OLAP, aggregates or denormalized tables will be computed.

This ETL job is costly to the production database. We may lose some data if inserts/updates are being processed during the transfer. Since we typically work with a lot of data, and we are storing aggregates, this is *usually* okay.

Addressing Performance: To improve performance on the RDBMS, we can

1. Have a replica database, and run an ETL job against the replica.
2. Dual write: on each write, write to the production DB as well as to the replica.

$$\begin{aligned} App &\longrightarrow Multiplexer^* \longrightarrow production\ DB \\ &\longrightarrow replica \xrightarrow{E} T \xrightarrow{L} DW \end{aligned}$$

*Multiplexer: Could be a load-balancer, message queue, etc.

If we have n grouping columns. we have $\sum_{k=0}^n \binom{n}{k} = 2^n$ dimensions.

In addition to aggregates, data warehouses can also store “exploded” (denormalized) tables for fast reads. This join can be computed as the data is inserted, and then quickly retrieved by the user without waiting. The join condition can be pre-specified in the schema.

OLAP v. OLTP

- OLAP cubes are computationally expensive.
- Reads of exploded tables are *fast*.
- Data is likely to be out of date.
- Typically append-only: no modifying/deleting.
- Read-only to internal users.

OLAP Operations

1. slice
2. dice
3. rollup
4. drill down
5. pivot

Slice: A slice selects one predominant dimension n from a cube and returns a new sub-cube, producing a rectangular representation of the data containing n . In SQL, we slice by `WHERE condition`.

Dice A dice selects multiple dimensions n from a cube and returns a new sub-cube, producing a rectangular representation of the data containing n . In SQL, we slice by doing `WHERE condition AND/OR ...`

Note: The distinction between slice and dice is usually only in theory. We usually slice.

Rollup A rollup computes aggregates across all levels of a hierarchical attribute. For example, $days \subseteq months \subseteq years \subseteq \dots$. In SQL, `ROLLUP(A1, ..., An)` where $A_n \subseteq \dots \subseteq A_1$. Also, `CUBE(A1, ..., An)` computes aggregates for all subsets of A_1, \dots, A_n .