

CS 143

Warren Kim

# Contents

0.1	Purpose of a Database . . . . .	2
0.2	Abstraction Layers . . . . .	2
0.3	Instances and Schema . . . . .	2
0.4	Data Models . . . . .	2
0.4.1	Relational . . . . .	3
0.4.2	Entity-Relationship (ER) . . . . .	3
0.4.3	Object-Oriented . . . . .	3
0.4.4	Document (Semi-Structured) . . . . .	3
0.4.5	Network/Hierarchical/Graphical . . . . .	3
0.4.6	Vector . . . . .	3
0.4.7	Key-Value . . . . .	3
0.5	Database Languages . . . . .	3
0.5.1	Data Manipulation Language . . . . .	3
0.5.2	Data Definition Language . . . . .	4
0.6	Data Storage and Querying . . . . .	4
0.7	Keys . . . . .	4
0.7.1	Superkey . . . . .	5
0.7.2	Candidate Key . . . . .	5
0.7.3	Primary Key . . . . .	5
0.7.4	Foreign Key . . . . .	5
0.8	Defining a Schema . . . . .	6
0.9	Relational Algebra . . . . .	6
0.9.1	Selection . . . . .	6
0.9.2	Projection . . . . .	6

## 0.1 Purpose of a Database

We will be studying (mostly) Relational DataBase Management Systems (RDBMS).

### Definition: Database

A **database** abstracts how data is stored, maintained, and processed. It is a system that uses advanced data structures to store and index data.

A database abstracts away the data integrity and file management aspect of CRUD operations. Moreover, a database provides us with a single location for all of the data, even if the database itself is distributed.

## 0.2 Abstraction Layers

There are three layers of abstraction: physical, logical, and view.

### Definition: Physical Abstraction

The **physical abstraction** defines the data and its relationships to other data within the database.

### Definition: Logical Abstraction

The **logical abstraction** deals with how we interface with the database.

### Definition: View Abstraction

The **view abstraction** refers to specific use cases and filters the data from the logical abstraction.

We start by learning the logical abstraction.

## 0.3 Instances and Schema

### Definition: Schema and Instance

A **schema**<sup>a</sup> is the overall design of a database. It defines the structure of the data as well as how it is organized.

An **instance** of a database is the actual set of data stored in the database at a particular moment in time.

---

<sup>a</sup>Note: schema can also refer to a relation (table).

## 0.4 Data Models

Data models define how we design databases and interact with data. We want to answer the following:

- (i) How do we define data?
- (ii) How do we encode relationships among data?
- (iii) How do we impose constraints on data?

Data models are either an Implementation model or a Design mechanism. Implementation models build databases from the ground up while design mechanisms are implemented as features in a database. We discuss five major types (an several niche ones).

### 0.4.1 Relational

In a relational model, all data is stored as a *relation*<sup>1</sup>. Rows represent individual  $n$ -tuple units (*records*). Columns represent (typed) *attributes* common to all records in the relations.

### 0.4.2 Entity-Relationship (ER)

An entity-relationship model uses a collection of basic objects (*entities*) and define *relationships* among them.

### 0.4.3 Object-Oriented

The object-oriented model is similar to OOP with encapsulation, methods, and object identity. It was originally an implementation model but is now a design mechanism.

### 0.4.4 Document (Semi-Structured)

A document model stores records as *documents*, which do *not* have an enforced schema. This allows for more versatility in the type of data stored in the database.

### 0.4.5 Network/Hierarchical/Graphical

A graph model is analogous to how we think. Records are stored as *nodes* and relationships between records as *edges*.

### 0.4.6 Vector

A vector model stores records as *vectors* in  $\mathbb{R}^n$ , and are stored in a way that enables efficient retrieval and comparison (e.g. nearest neighbor[s]).

### 0.4.7 Key-Value

A key-value model stores data as a key-value pair (typically using a hash function). In this model, data typically lives in RAM as opposed to disk.

## 0.5 Database Languages

There are two main semantic systems when working with databases:

- (i) Data Manipulation Language (DML)
- (ii) Data Definition Language (DDL)

Note that a relational model typically uses SQL for both DDL and DML.

### 0.5.1 Data Manipulation Language

DML's can either be procedural or declarative.

Definition: Query

A *query* is a written expression to retrieve or manipulate data.

---

<sup>1</sup>Note: tables are an implementation of relations.

### Aside: A Note on SQL

SQL is a declarative language, and as such, it is hard to perform sequential or nontrivial<sup>a</sup> computations in SQL. To remedy this, a common option is to write an **ETL job** in another language (pick one). We **E**xtract the data from the database (using a connection driver), **T**ransform the data using another language (pick one!), and **L**oad the data into a new table using the same driver. We can schedule these jobs using something like **cron**.

<sup>a</sup>Nontrivial: Any computation where we have to specify *how* to perform the computation.

## 0.5.2 Data Definition Language

DDL's specify a schema: a collection of attribute names and data types, consistency constraints, and optionally storage structure and access methods. There are four types of consistency constraints:

- (i) Domain constraints define the domain of an attribute (e.g. **tinyint**, **enum**, etc.).
- (ii) Assertions are business rules that must hold true (e.g. an enforced prerequisite for a class must be present in your transcript before you can add a class to your study list).
- (iii) Authorization determines who can do what (e.g. full CRUD, read-only, etc.).
- (iv) Referential integrity ensures that links from one table to another must be defined (Suppose we have two relations  $R, R'$ . If there is a link  $f : R \rightarrow R'$ , then  $f$  is surjective).

## 0.6 Data Storage and Querying

### Definition: Storage Manager

A **storage manager** that abstracts away how the data is laid out on disk.

A storage manager is helpful because reading data from disk to RAM is *slow*, and the storage manager handles swapping<sup>2</sup> and makes retrieval efficient.

### Definition: Query Manager

A **query manager** takes the DML statements and organizes them into a *query plan*<sup>a</sup> that “compiles” a query (using relational algebra) and executes the instruction(s).

<sup>a</sup>Note: The query plan dictates the performance of a query.

## 0.7 Keys

### Aside: A Note on Context and Instance

Based on **context** means that the given data is a subset of the complete dataset.  
Based on **instance** means that we treat the given data as the complete dataset.

<sup>2</sup>Swapping: Virtual memory in CS111!

### 0.7.1 Superkey

#### Definition: Superkey

A **superkey** is a set of one or more attributes that uniquely identifies a record (tuple) and distinguishes it from all other records in the relation.

Formally, let  $R$  be a relation with a set  $S = \{a_1, a_2, \dots, a_n : a \text{ is an attribute of } R\}$ . A **superkey** is a subset  $s \subseteq S$  such that  $s$  uniquely identifies each  $n$ -tuple in  $R$ .

The superkey  $s = S = \{a_1, a_2, \dots, a_n\} = \bigcup_{i=1}^n \{a_i\}$  is called the **trivial superkey**. Additionally,  $\emptyset$  is not a superkey. Further note that for every relation  $R$ , there exists at most  $2^n - 1$  superkeys where  $n$  is the number of attributes.

### 0.7.2 Candidate Key

#### Definition: Candidate Key

A **candidate key** is a superkey such that no subset of the candidate key is a superkey; i.e. it is the minimal superkey.

Formally, let  $R$  be a relation with a set  $S = \{a_1, a_2, \dots, a_n : a \text{ is an attribute of } R\}$ . A **candidate key** is a superkey  $s \subseteq S$  such that for every proper subset  $t \subsetneq s$ ,  $t$  is not a superkey.

Candidate keys may vary in length, and the attributes of a candidate key may be NULL as long as it uniquely identifies an  $n$ -tuple in the relation.

### 0.7.3 Primary Key

#### Definition: Primary Key and Composite Key

A **primary key** is a candidate key (chosen by the database designer) to enforce uniqueness for a particular use case.

The primary key is typically chosen to be the minimal candidate key for simplicity. The attributes of a primary key may not be NULL.

### 0.7.4 Foreign Key

#### Definition: Foreign Key

A **foreign key** is a set of attributes that links tuples of two relations.

Formally, let  $R, R'$  be relations with sets  $S = \{a_1, a_2, \dots, a_n : a \text{ is an attribute of } R\}, S' = \{a'_1, a'_2, \dots, a'_n : a' \text{ is an attribute of } R'\}$ . A **foreign key** is a key  $s \subseteq S$  of  $R$  that maps to the primary key  $p \subseteq S'$  of  $R'$ .

Foreign keys are used to enforce referential integrity constraints; i.e. foreign keys in a relation  $R$  are used to protect data in  $R$  from being orphaned and/or inconsistent. Given two relations  $R, R'$  related via a foreign key,  $R'$  is said to be the *referring* relation and  $R$  the *referred* relation.

Let two relations  $R, S$  be related via a foreign key, where  $S$  is the *referring* relation and  $R$  is the *referred* relation. Suppose we want to remove an  $n$ -tuple  $r \in R$ . Then there are two cases:

*Case 1* If there is no  $s \in S$  such that  $s \mapsto r$ , we simply remove  $r$ .

*Case 2* If there is at least one  $s \in S$  such that  $s \mapsto r$ , we can either throw an error to prevent the deletion of  $r$  or *cascade*<sup>3</sup> the delete.

<sup>3</sup>Cascade: Delete  $r$  and all  $s \in S$  that refer to  $r$ .

## 0.8 Defining a Schema

A schema can be written as `relation(attribute1, ..., attributen)` where underlined attributes represent the primary key.

## 0.9 Relational Algebra

describe relational algebra

### 0.9.1 Selection

Definition: Selection

**Selection** retrieves a subset of tuples from a *single* relation  $R$  that satisfies some predicate  $\psi$  and returns a new relation  $R' \subseteq R$ , and is defined by

$$\sigma_{\psi}(R) = R' = \{t \in R : \psi(t)\}$$

where  $\psi$  is a boolean predicate on attributes and values with respect to a unary or binary operator<sup>a</sup>

---

<sup>a</sup>We may use the following operators:  $\{=, \neq, <, >, \leq, \geq, \neg\}$ .

We can build complex predicates using conjunction  $\wedge$  (and) or disjunction  $\vee$  (or).

**Note: that selection  $\sigma$  is implemented as WHERE in SQL.**

Below are a list of examples of selection, assuming all attributes and relations are well-defined:

- (i)  $\sigma_{(\text{dislikes} < \text{likes})}(\text{youtube\_video})$
- (ii)  $\sigma_{(\text{cat\_id}=17)}(\text{youtube\_video})$
- (iii)  $\sigma_{([\text{dislikes} < \text{likes}] \wedge [\text{views} > 1000000] \wedge [\text{cat\_id}=24])}(\text{youtube\_video})$
- (iv)  $\sigma_{(\text{dislikes} < \text{likes})}(\sigma_{(\text{views} > 1000000)}(\sigma_{(\text{cat\_id}=24)}(\text{youtube\_videos})))$

Note that (iii) and (iv) are equivalent.

### 0.9.2 Projection

Definition: Projection

**Projection** extracts attributes from a set of tuples and removes duplicates. Given a relation  $R$ ,  $n$ -tuple  $t$ , and a set of attributes  $a_1, \dots, a_n$ ,

$$\Pi_{a_1, \dots, a_n}(R) = \{t[a_1, \dots, a_n] : t \in R\}$$

Projection is usually the last (outermost) operation done on a relation.

Aside: Projection?

We call it a projection because we are collapsing an  $n$ -tuple down to an  $(n - k)$ -tuple. That is, we take the  $n$ -tuples in a relation  $R_n$  and collapse them into a set of  $(n - k)$ -tuples in a new relation  $R'_{n-k}$ .

---

Here,  $R_n$  is a relation with  $n$  attributes.